


AudioMNIST controls at 24,000 steps



Hollow: individual training seeds. Filled: mean $\pm$ population SD. Upper-left is better.

* Original-checkpoint reevaluation differs from archive; new gain points are measured. Historical points remain unchanged; baseline reproduction is not established.